

1 K-means

1.1 What's the algorithm?

The K-means algorithm is an unsupervised learning algorithm that divides data points into K clusters. The algorithm consists of two parts, where the last part is run as many times as we, the user, specifies, or until the algorithm has converged:

1. Initialize **cluster centroids** $\mu_1, \mu_2, \dots, \mu_k \in \mathbb{R}^d$ randomly.
2. Repeat:
 - (a) For every data point i , set $c^{(i)} := \arg \min_j ||x^{(i)} - \mu_j||^2$, as the centroid closest to point i .
 - (b) For each centroid j , set $\mu_{j,new} := \frac{\sum_{i=1}^n \mathbb{1}\{c^{(i)}=j\}x^{(i)}}{\sum_{i=1}^n \mathbb{1}\{c^{(i)}=j\}}$, as the new position for centroid j .

1.2 Inductive bias

The inductive bias in this problem is that we assume that there are K clusters that we want to find, and that the clusters are of similar size.

In the second dataset, the ratios of the x_0 's and x_1 's were different than what was given in the first dataset. To fix this I divided the x_0 's by 10 to get the ratios equal to one.

1.3 Some plots

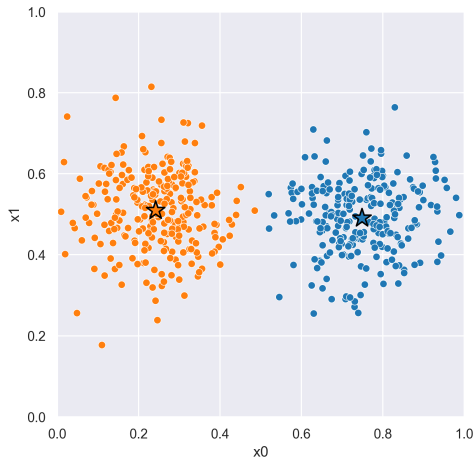


Figure 1: Clusters in dataset 1

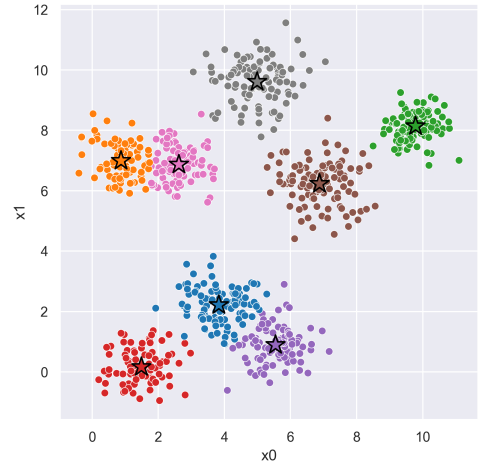


Figure 2: Clusters in dataset 2

1.4 Result/Preprocessing

For the first dataset I didn't have to do any preprocessing, but I matched the result given as an example. In the second dataset, the ratios of the x 's were a lot different than one, so to fix this I divided the x_0 's by 10, to get the desired result as seen above, where all 8 clusters are found.

2 Logistic regression

2.1 What's the algorithm?

Logistic regression works by fitting $\mathbf{u} = (\mathbf{w}, b)$ such that

$$J(\mathbf{u}) = \sum_{i=1}^N (-y_i \log(\sigma_i) - (1 - y_i) \log(1 - \sigma_i))$$

is minimized. Here $\sigma_i = \sigma(\hat{y}_i)$. This gets done using gradient descent:

1. Initialize \mathbf{u} randomly.
2. Repeat:
 - (a) Calculate the gradient $\nabla = \frac{\partial J}{\partial \mathbf{u}}$
 - (b) Update $\mathbf{u} = \mathbf{u} - \eta \nabla$, where η is the gradient step size.
3. until $\frac{\|\mathbf{u} - \mathbf{u}^{old}\|}{\|\mathbf{u}^{old}\|} < \epsilon$.

2.2 Inductive bias

The inductive bias in this problem is that we assume that the data has two different classes, i.e. $y = 1$ or $y = 0$, and we assume that there is a linear boundary between these two classes.

The second dataset had the two classes separated in a circle, so to get around this problem I used the distance from the mean position of all the points to separate the two.

2.3 Some plots

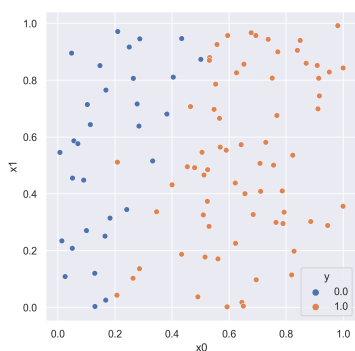


Figure 3

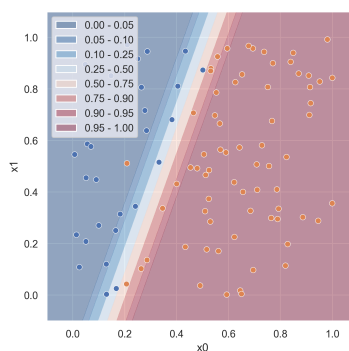


Figure 4

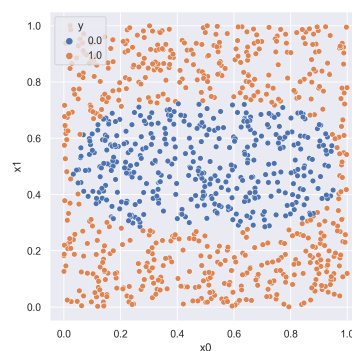


Figure 5

2.4 Result/Preprocessing

In the first dataset I didnt have to do any preprocessing, but I exceeded the accuracy and cross entropy in the example. I got an accuracy of 0.980 and a cross entropy of 0.097. In the second dataset, there wasnt a linear barrier between the groups, so I used the distance from the middle to make such a barrier. The training data then gave an accuracy of 0.818 and a cross entropy of 0.365, and the testing data gave an accuracy of 0.826 and a cross entropy of 0.332.